Ugrid Reference

Contents

1 Overview

Ugrid is a fast and general purpose distributed data processing system. It provides a high-level API in Javascript and an optimized parallel execution engine.

A Ugrid application consists of a *master* program that runs the user code and executes various *parallel operations* on a cluster of *workers*.

The main abstraction Ugrid provides is a *dataset* which is similar to a Javascript *array*, but partitioned accross the workers that can be operated in parallel.

There are several ways to create a dataset: *parallelizing* an existing array in the master program, or referencing a dataset in a distributed storage system (such as HDFS), or *streaming* the content of any source that can be processed through Node.js *Streams*. We call *source* a function which initializes a dataset.

Datasets support two kinds of operations: *transformations*, which create a new dataset from an existing one, and *actions*, which return a value to the *master* program after running a computation on the dataset.

For example, `map` is a transformation that applies a function to each element of a dataset, returning a new dataset. On the other hand, `reduce` is an action that aggregates all elements of a dataset using some function, and returns the final result to the master.

Sources and *transformations* in Ugrid are *lazy*. They do not start right away, but are triggered by *actions*, thus allowing efficient pipelined execution and optimized data transfers.

A first example:

```javascript
var uc = require('ugrid').context();      // create a new context
uc.parallelize([1, 2, 3, 4]).            // source
  map(function (x) {return x+1}).         // transform
  reduce(function (a, b) {return a+b}, 0). // action
  then(console.log);                       // process result: 14
```

2 Working with datasets

After having initialized a cluster context using ugrid.context(), one can create a dataset using the following sources:

Source Name	Description
lineStream(stream)	Create a dataset from a text stream
objectStream(stream)	Create a dataset from an object stream
parallelize(array)	Create a dataset from an array
textFile(path)	Create a dataset from a regular text file

Transformations operate on a dataset and return a new dataset. Note that some transformation operate only on datasets where each element is in the form of 2 elements array of key and value (`[k,v]` dataset):

```
[[Ki,Vi], ..., [Kj, Vj]]
```

A special transformation `persist()` enables one to *persist* a dataset in memory, allowing efficient reuse across parallel operations.

Transformation Name	Description	in	out
cartesian(other)	Perform a cartesian product with the other dataset	v w	[v,w]

Transformation Name	Description	in	out
coGroup(other)	Group data from both datasets sharing the same key	[k,v] [k,w]	[k,[[v],[w]]]
distinct()	Return a dataset where duplicates are removed	v	w
filter(func)	Return a dataset of elements on which function returns true	v	w
flatMap(func)	Pass the dataset elements to a function which returns a sequence	v	w
groupByKey()	Group values with the same key	[k,v]	[k,[v]]
intersection(other)	Return a dataset containing only elements found in both datasets	v w	v
join(other)	Perform an inner join between 2 datasets	[k,v]	[k,[v,w]]
leftOuterJoin(other)	Join 2 datasets where the key must be present in the other	[k,v]	[k,[v,w]]
rightOuterJoin(other)	Join 2 datasets where the key must be present in the first	[k,v]	[k,[v,w]]
keys()	Return a dataset of just the keys	[k,v]	k
map(func)	Return a dataset where elements are passed through a function	v	w
mapValues(func)	Map a function to the value field of key-value dataset	[k,v]	[k,w]
reduceByKey(func, init)	Combine values with the same key	[k,v]	[k,w]
persist()	Idempotent. Keep content of dataset in cache for further reuse.	v	v
sample(rep, frac, seed)	Sample a dataset, with or without replacement	v	w
subtract(other)	Remove the content of one dataset	v w	v
union(other)	Return a dataset containing elements from both datasets	v	v w
values()	Return a dataset of just the values	[k,v]	v

Actions operate on a dataset and return a result to the *master*. Results are always returned asynchronously. In a case of a single result, it is returned through a either a callback, or an ES6 promise. In the case of multiple value results, the action returns a readable stream.

Action Name	Description	out
aggregate(func, func, init)	Similar to reduce() but may return a different type	value
collect()	Return the content of dataset	stream of elements
count()	Return the number of elements from dataset	number
countByKey()	Return the number of occurrences for each key in a `[k,v]` dataset	stream of [k,number]
countByValue()	Return the number of occurrences of elements from dataset	stream of [v,number]
foreach(func)	Apply the provided function to each element of the dataset	**not implemented**
lookup(k)	Return the list of values `v` for key `k` in a `[k,v]` dataset	stream of v
reduce(func, init)	Aggregates dataset elements using a function, return a single value	value

3   Ugrid module

The Ugrid module is the main entry point for Ugrid functionality. To use it, one must `require('ugrid')`.

3.1   ugrid.context([config])

Creates and returns a new context which represents the connection to the Ugrid cluster, and which can be used to create datasets on that cluster. Config is an *Object* which defines the cluster server, with the following defaults:

```
{
  host: 'localhost',    // Cluster server host, settable also by UGRID_HOST env
  port: '12346'         // Cluster server port, settable also by UGRID_PORT env
}
```

Example:

```
var ugrid = require('ugrid');
var uc = ugrid.context();
```

3.1.1   uc.end()

Closes the connection to the cluster.

3.1.2   uc.parallelize(array)

Returns a new dataset containing elements from the *Array* array.

Example:

```
var a = uc.parallelize(['Hello', 'World']);
```

3.1.3   uc.textFile(path)

Returns a dataset of lines composing the file specified by path *String*.

Note: If using a path on the local filesystem, the file must also be accessible at the same path on worker nodes. Either copy the file to all workers or use a network-mounted shared file system.

Example, the following program prints the length of a text file:

```
var lines = uc.textFile('data.txt');
lines.map(s => s.length).reduce((a, b) => a + b, 0).then(console.log);
```

3.1.4   uc.lineStream(input_stream)

Returns a dataset of lines of text read from input_stream *Object*, which is a readable stream where dataset content is read from.

The following example computes the size of a file using streams:

```
var stream = fs.createReadStream('data.txt', 'utf8');
uc.lineStream(stream).
   map(s => s.length).
```

```
  reduce((a, b) => a + b, 0).
  then(console.log);
```

3.1.5 uc.objectStream(input_stream)

Returns a dataset of Javascript *Objects* read from input_stream *Object*, which is a readable stream where dataset content is read from.

The following example counts the number of objects returned in an object stream using the mongodb native Javascript driver:

```
var cursor = db.collection('clients').find();
uc.objectStream(cursor).count().then(console.log);
```

3.2 Dataset methods

Dataset objects, as created initially by above ugrid context source functions, have the following methods, allowing either to instantiate a new dataset through a transformation, or to return results to the master program.

3.2.1 ds.aggregate(reducer, combiner, init[,obj][,done])

Returns the aggregated value of the elements of the dataset using two functions *reducer()* and *combiner()*, allowing to use an arbitrary accumulator type, different from element type (as opposed to `reduce()` which imposes the same type for accumulator and element). The result is passed to the *done()* callback if provided, otherwise an ES6 promise is returned.

- *reducer*: a function of the form `function(acc,val[,obj[,wc]])`, which returns the next value of the accumulator (which must be of the same type as *acc*) and with:
  - *acc*: the value of the accumulator, initially set to *init*
  - *val*: the value of the next element of the dataset on which `aggregate()` operates
  - *obj*: the same parameter *obj* passed to `aggregate()`
  - *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies.
- *combiner*: a function of the form `function(acc1,acc2[,obj])`, which returns the merged value of accumulators and with:
  - *acc1*: the value of an accumulator, computed locally on a worker
  - *acc2*: the value of an other accumulator, issued by another worker
  - *obj*: the same parameter *obj* passed to `aggregate()`
- *init*: the initial value of the accumulators that are used by *reducer()* and *combiner()*. It should be the identity element of the operation (i.e. applying it through the function should not change result).
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset
- *done*: a callback of the form `function (error, result)` which is called at completion. If *undefined*, `aggregate()` returns an ES6 promise.

The following example computes the average of a dataset, avoiding a `map()`:

```
uc.parallelize([3, 5, 2, 7, 4, 8]).
   aggregate((a, v) => [a[0] + v, a[1] + 1],
       (a1, a2) => [a1[0] + a2[0], a1[1] + a2[1]],
       [0, 0],
       function (err, res) {
           console.log(res[0] / res[1]);
```

```
        });
// 4.8333
```

3.2.2   ds.cartesian(other)

Returns a dataset wich contains all possible pairs `[a, b]` where `a` is in the source dataset and `b` is in the *other* dataset.

Example:

```
var ds1 = uc.parallelize([1, 2]);
var ds2 = uc.parallelize(['a', 'b', 'c']);
ds1.cartesian(ds2).collect().toArray().then(console.log);
// [ [ 1, 'a' ], [ 1, 'b' ], [ 1, 'c' ],
//   [ 2, 'a' ], [ 2, 'b' ], [ 2, 'c' ] ]
```

3.2.3   ds.coGroup(other)

When called on dataset of type `[k,v]` and `[k,w]`, returns a dataset of type `[k, [[v], [w]]]`, where data of both datasets share the same key.

Example:

```
var ds1 = uc.parallelize([[10, 1], [20, 2]]);
var ds2 = uc.parallelize([[10, 'world'], [30, 3]]);
ds1.coGroup(ds2).collect().on('data', console.log);
// [ 10, [ [ 1 ], [ 'world' ] ] ]
// [ 20, [ [ 2 ], [] ] ]
// [ 30, [ [], [ 3 ] ] ]
```

3.2.4   ds.collect([opt])

Returns a readable stream of all elements of the dataset. Optional *opt* parameter is an object with the default content `{text: false}`. if `text` option is `true`, each element is passed through `JSON.stringify()` and a 'newline' is appended, making it possible to pipe to standard output or any text stream.

Example:

```
uc.parallelize([1, 2, 3, 4]).
   collect({text: true}).pipe(process.stdout);
// 1
// 2
// 3
// 4
```

3.2.5   ds.count([callback])

Returns the number of elements in the dataset.

The *callback* is the form of `function(error, result)`, and is called asynchronously at completion. In *undefined*, an ES6 promise is returned.

Example:

```
uc.parallelize([10, 20, 30, 40]).count().then(console.log);
// 4
```

3.2.6   ds.countByKey()

When called on a dataset of type `[k,v]`, computes the number of occurrences of elements for each key in a dataset of type `[k,v]`. Returns a readable stream of elements of type `[k,w]` where `w` is the result count.

Example:

```
uc.parallelize([[10, 1], [20, 2], [10, 4]]).
   countByKey().on('data', console.log);
// [ 10, 2 ]
// [ 20, 1 ]
```

3.2.7   ds.countByValue()

Computes the number of occurences of each element in dataset and returns a readable stream of elements of type `[v,n]` where `v` is the element and `n` its number of occurrences.

Example:

```
uc.parallelize([ 1, 2, 3, 1, 3, 2, 5 ]).
   countByValue().
   toArray().then(console.log);
// [ [ 1, 2 ], [ 2, 2 ], [ 3, 2 ], [ 5, 1 ] ]
```

3.2.8   ds.distinct()

Returns a dataset where duplicates are removed.

Example:

```
uc.parallelize([ 1, 2, 3, 1, 4, 3, 5 ]).
   distinct().
   collect().toArray().then(console.log);
// [ 1, 2, 3, 4, 5 ]
```

3.2.9   ds.filter(filter[,obj])

- *filter*: a function of the form `callback(element[,obj[,wc]])`, returning a *Boolean* and where:
  - *element*: the next element of the dataset on which `filter()` operates
  - *obj*: the same parameter *obj* passed to `filter()`
  - *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies.
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset

Applies the provided filter function to each element of the source dataset and returns a new dataset containing the elements that passed the test.

Example:

```
function filter(data, obj) { return data % obj.modulo; }

uc.parallelize([1, 2, 3, 4]).
   filter(filter, {modulo: 2}).
   collect().on('data', console.log);
// 1 3
```

3.2.10 ds.flatMap(flatMapper[,obj])

Applies the provided mapper function to each element of the source dataset and returns a new dataset.

- *flatMapper*: a function of the form `callback(element[,obj[,wc]])`, returning an *Array* and where:
    - *element*: the next element of the dataset on which `flatMap()` operates
    - *obj*: the same parameter *obj* passed to `flatMap()`
    - *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies.
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset

Example:

```
function flatMapper(data, obj) {
    var tmp = [];
    for (var i = 0; i < obj.N; i++) tmp.push(data);
    return tmp;
}

uc.parallelize([1, 2, 3, 4]).
   flatMap(flatMapper, {N: 2}).
   collect().on('data', console.log);
// [ 'hello', 2 ]
// [ 'hello', 2 ]
// [ 'world', 4 ]
// [ 'world', 4 ]
```

3.2.11 ds.flatMapValues(flatMapper[,obj])

Applies the provided flatMapper function to the value of each [key, value] element of the source dataset and return a new dataset containing elements defined as [key, mapper(value)], keeping the key unchanged for each source element.

- *flatMapper*: a function of the form `callback(element[,obj[,wc]])`, returning an *Array* and where:
    - *element*: the value v of the next [k,v] element of the dataset on which `flatMapValues()` operates
    - *obj*: the same parameter *obj* passed to `flatMapValues()`
    - *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies.
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset

Example:

```
function valueFlatMapper(data, obj) {
    var tmp = [];
    for (var i = 0; i < obj.N; i++) tmp.push(data * obj.fact);
    return tmp;
}

uc.parallelize([['hello', 1], ['world', 2]]).
   flatMapValues(valueFlatMapper, {N: 2, fact: 2}).
   collect().on('data', console.log);
```

3.2.12   ds.foreach(callback[, obj][, done])

not implemented

This action applies a *callback* function on each element of the dataset.

- *callback*: a function of the form `function(val[,obj[,wc]])`, which returns *null* and with:
  - *val*: the value of the next element of the dataset on which `foreach()` operates
  - *obj*: the same parameter *obj* passed to `foreach()`
  - *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies.
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset
- *done*: a callback of the form `function (error, result)` which is called at completion. If *undefined*, `foreach()` returns an ES6 promise.

In the following example, the `console.log()` callback provided to `foreach()` is executed on workers and may be not visible:

```
uc.parallelize([1, 2, 3, 4]).
   foreach(console.log).then(console.log('finished'));
```

3.2.13   ds.groupByKey()

When called on a dataset of type `[k,v]`, returns a dataset of type `[k, [v]]` where values with the same key are grouped.

Example:

```
uc.parallelize([[10, 1], [20, 2], [10, 4]]).
   groupByKey().collect().on('data', console.log);
// [ 10, [ 1, 4 ] ]
// [ 20, [ 2 ] ]
```

3.2.14   ds.intersection(other)

Returns a dataset containing only elements found in source dataset and *other* dataset.

Example:

```
var ds1 = uc.parallelize([1, 2, 3, 4, 5]);
var ds2 = uc.parallelize([3, 4, 5, 6, 7]);
ds1.intersection(ds2).collect().toArray().then(console.log);
// [ 3, 4, 5 ]
```

3.2.15   ds.join(other)

When called on source dataset of type `[k,v]` and *other* dataset of type `[k,w]`, returns a dataset of type `[k, [v, w]]` pairs with all pairs of elements for each key.

Example:

```
var ds1 = uc.parallelize([[10, 1], [20, 2]]);
var ds2 = uc.parallelize([[10, 'world'], [30, 3]]);
ds1.join(ds2).collect().on('data', console.log);
// [ 10, [ 1, 'world' ] ]
```

3.2.16   ds.keys()

When called on source dataset of type `[k,v]`, returns a dataset with just the elements `k`.

Example:

```
uc.parallelize([[10, 'world'], [30, 3]]).
   keys.collect().on('data', console.log);
// 10
// 30
```

3.2.17   ds.leftOuterJoin(other)

When called on source dataset of type `[k,v]` and *other* dataset of type `[k,w]`, returns a dataset of type `[k, [v, w]]` pairs where the key must be present in the *other* dataset.

Example:

```
var ds1 = uc.parallelize([[10, 1], [20, 2]]);
var ds2 = uc.parallelize([[10, 'world'], [30, 3]]);
ds1.leftOuterJoin(ds2).collect().on('data', console.log);
// [ 10, [ 1, 'world' ] ]
// [ 20, [ 2, null ] ]
```

3.2.18   ds.lookup(k)

When called on source dataset of type `[k,v]`, returns a readable stream of values `v` for key `k`.

Example:

```
uc.parallelize([[10, 'world'], [20, 2], [10, 1], [30, 3]]).
   lookup(10).on('data', console.log);
// world
// 1
```

3.2.19   ds.map(mapper[,obj])

Applies the provided mapper function to each element of the source dataset and returns a new dataset.

- *mapper*: a function of the form `callback(element[,obj[,wc]])`, returning an element and where:
  - *element*: the next element of the dataset on which `map()` operates
  - *obj*: the same parameter *obj* passed to `map()`
  - *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies.
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset

The following example program

```
uc.parallelize([1, 2, 3, 4]).
   map((data, obj) => data * obj.scaling, {scaling: 1.2}).
   collect().toArray().then(console.log);
// [ 1.2, 2.4, 3.6, 4.8 ]
```

3.2.20   ds.mapValues(mapper[,obj])

- *mapper*: a function of the form `callback(element[,obj[,wc]])`, returning an element and where:
  - *element*: the value v of the next [k,v] element of the dataset on which `mapValues()` operates
  - *obj*: the same parameter *obj* passed to `mapValues()`
  - *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset

Applies the provided mapper function to the value of each `[k,v]` element of the source dataset and return a new dataset containing elements defined as `[k, mapper(v)]`, keeping the key unchanged for each source element.

Example:

```
uc.parallelize([['hello', 1], ['world', 2]]).
   mapValues((a, obj) => a*obj.fact, {fact: 2}).
   collect().on('data', console.log);
// ['hello', 2]
// ['world', 4]
```

3.2.21   ds.reduce(reducer, init[,obj][,done])

Returns the aggregated value of the elements of the dataset using a *reducer()* function. The result is passed to the *done()* callback if provided, otherwise an ES6 promise is returned.

- *reducer*: a function of the form `function(acc,val[,obj[,wc]])`, which returns the next value of the accumulator (which must be of the same type as *acc* and *val*) and with:
  - *acc*: the value of the accumulator, initially set to *init*
  - *val*: the value of the next element of the dataset on which `reduce()` operates
  - *obj*: the same parameter *obj* passed to `reduce()`
  - *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies.
- *init*: the initial value of the accumulators that are used by *reducer()*. It should be the identity element of the operation (i.e. applying it through the function should not change result).
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset
- *done*: a callback of the form `function (error, result)` which is called at completion. If *undefined*, `reduce()` returns an ES6 promise.

Example:

```
uc.parallelize([1, 2, 4, 8]).
   reduce((a, b) => a + b, 0).
   then(console.log);
// 15
```

3.2.22   ds.reduceByKey(reducer, init[, obj])

- *reducer*: a function of the form `callback(acc,val[,obj[,wc]])`, returning the next value of the accumulator (which must be of the same type as *acc* and *val*) and where:
  - *acc*: the value of the accumulator, initially set to *init*
  - *val*: the value v of the next [k,v] element of the dataset on which `reduceByKey()` operates
  - *obj*: the same parameter *obj* passed to `reduceByKey()`

- *wc*: the worker context, a persistent object local to each worker, where user can store and access worker local dependencies.
- *init*: the initial value of accumulator for each key. Will be passed to *reducer*.
- *obj*: user provided data. Data will be passed to carrying serializable data from master to workers, obj is shared amongst mapper executions over each element of the dataset

When called on a dataset of type `[k,v]`, returns a dataset of type `[k,v]` where the values of each key are aggregated using the *reducer* function and the *init* initial value.

Example:

```
uc.parallelize([[10, 1], [10, 2], [10, 4]]).
   reduceByKey((a,b) => a+b, 0).
   collect().on('data', console.log);
// [10, 7]
```

3.2.23   ds.rightOuterJoin(other)

When called on source dataset of type `[k,v]` and *other* dataset of type `[k,w]`, returns a dataset of type `[k, [v, w]]` pairs where the key must be present in the *source* dataset.

Example:

```
var ds1 = uc.parallelize([[10, 1], [20, 2]]);
var ds2 = uc.parallelize([[10, 'world'], [30, 3]]);
ds1.rightOuterJoin(ds2).collect().on('data', console.log);
// [ 10, [ 1, 'world' ] ]
// [ 30, [ null, 2 ] ]
```

3.2.24   ds.sample(withReplacement, frac, seed)

- *withReplacement*: *Boolean* value, *true* if data must be sampled with replacement
- *frac*: *Number* value of the fraction of source dataset to return
- *seed*: *Number* value of pseudo-random seed

Returns a dataset by sampling a fraction *frac* of source dataset, with or without replacement, using a given random generator *seed*.

Example:

```
uc.parallelize([1, 2, 3, 4, 5, 6, 7, 8]).
   sample(true, 0.5, 0).
   collect().toArray().then(console.log);
// [ 1, 1, 3, 4, 4, 5, 7 ]
```

3.2.25   ds.subtract(other)

Returns a dataset containing only elements of source dataset which are not in *other* dataset.

Example:

```
var ds1 = uc.parallelize([1, 2, 3, 4, 5]);
var ds2 = uc.parallelize([3, 4, 5, 6, 7]);
ds1.subtract(ds2).collect().on('data', console.log);
// 1 2
```

3.2.26   ds.union(other)

Returns a dataset that contains the union of the elements in the source dataset and the *other* dataset.

Example:

```
var ds1 = uc.parallelize([1, 2, 3, 4, 5]);
var ds2 = uc.parallelize([3, 4, 5, 6, 7]);
ds1.union(ds2).collect().toArray().then(console.log);
// [ 1, 2, 3, 4, 5, 3, 4, 5, 6, 7 ]
```

3.2.27   ds.values()

When called on source dataset of type `[k,v]`, returns a dataset with just the elements `v`.

Example:

```
uc.parallelize([[10, 'world'], [30, 3]]).
   keys.collect().on('data', console.log);
// 'world'
// 3
```

4   References

ES6 promise https://promisesaplus.com

readable stream https://nodejs.org/api/stream.html#stream_class_stream_readable